

Cloud ML Engine

Google Cloud Platform (GCP)

2017.08

박찬성

before getting started..

1. do you have a GCP account? if no, just create one
 1. if you have a gmail account, you are ok to go
2. just read THIS, it will give a basic idea how everything works. You don't need to follow every instructions. just reading is sufficient
3. do some of the pre-requisite steps introduced like..

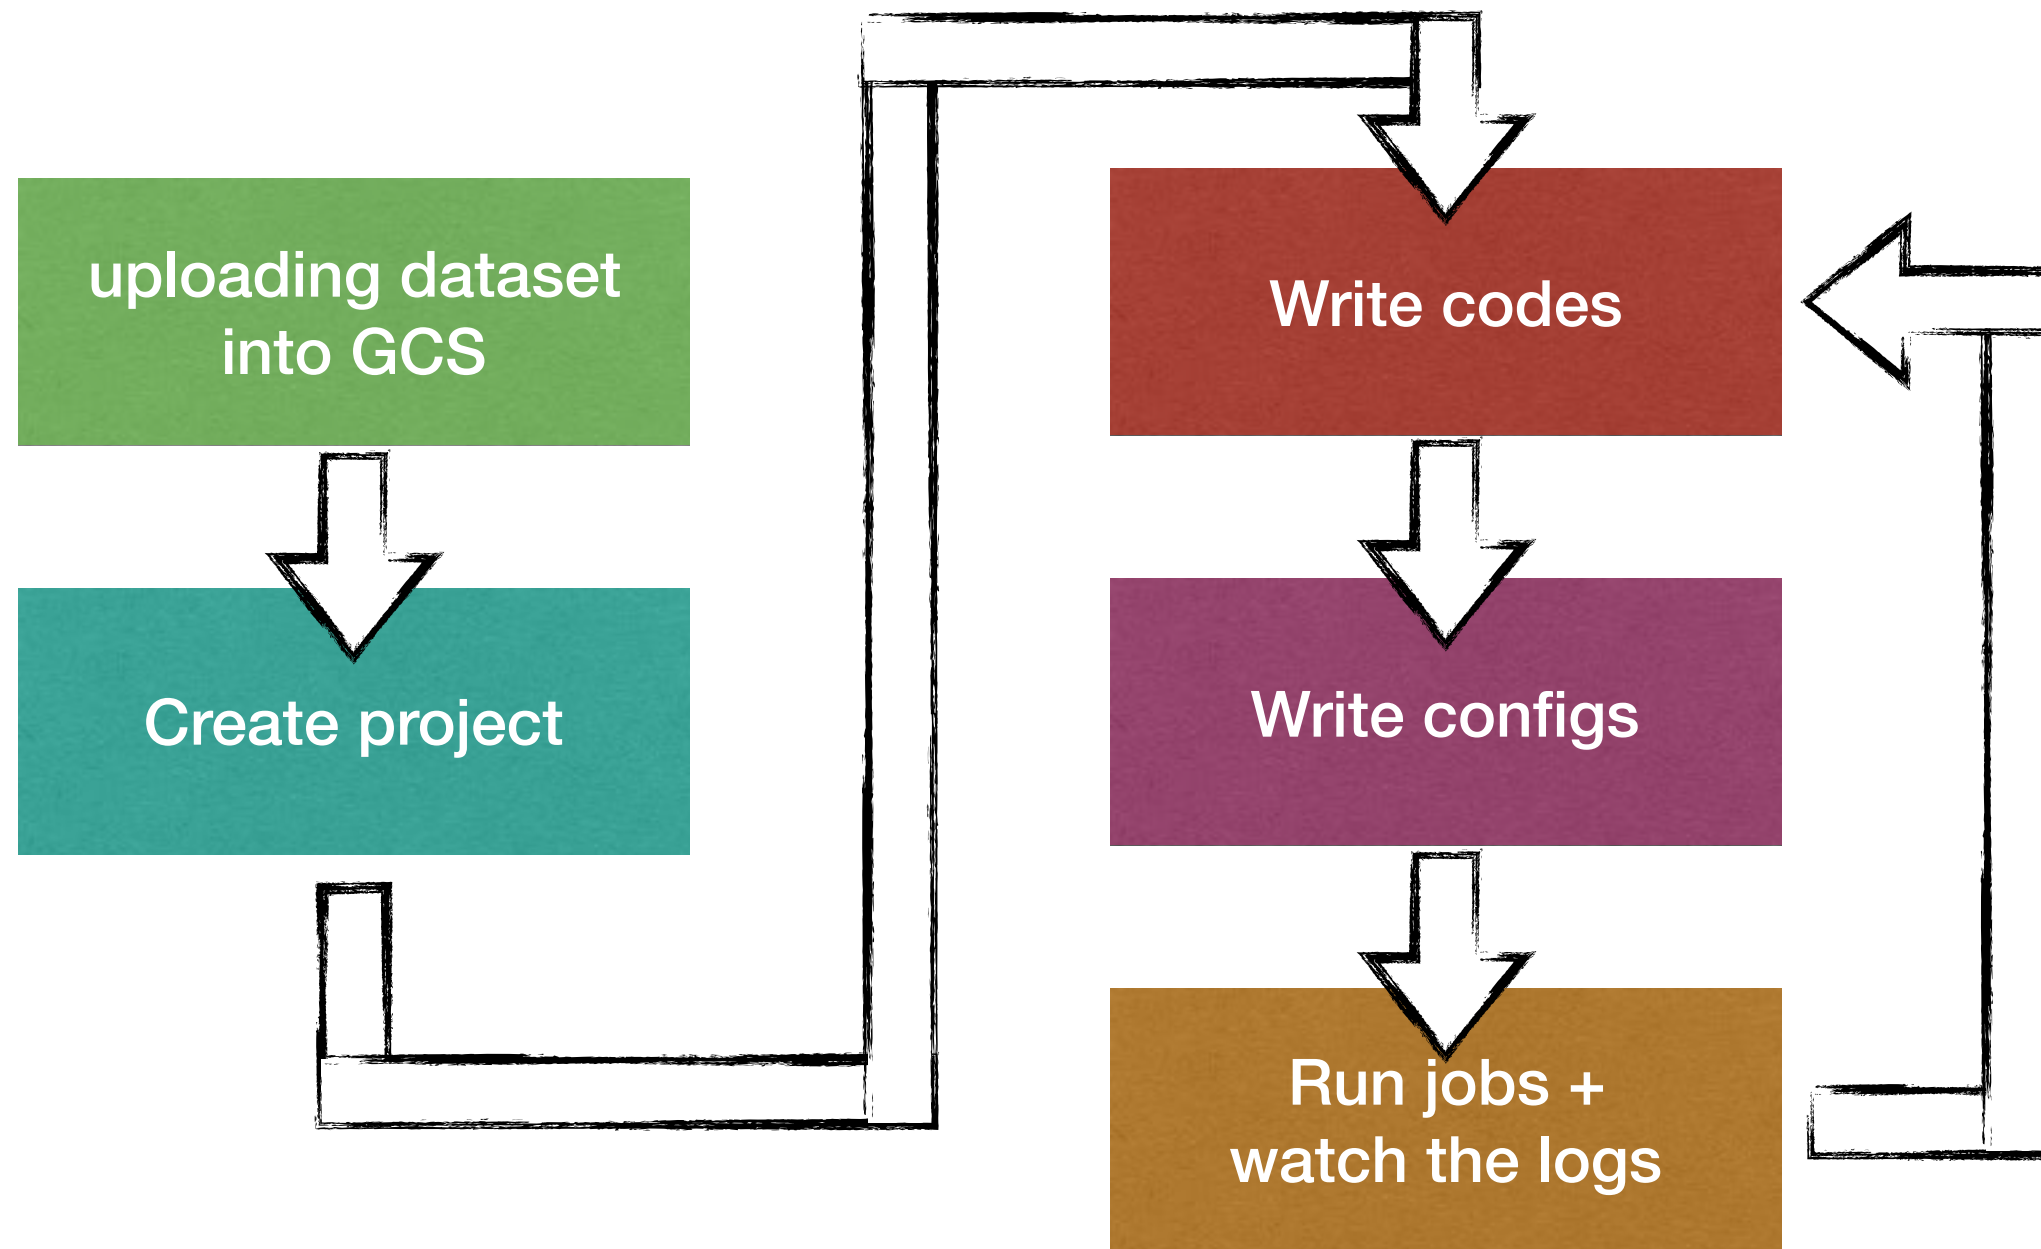
The screenshot shows the 'Set up your GCP project' page. It contains a list of steps with corresponding buttons and examples.

Set up your GCP project

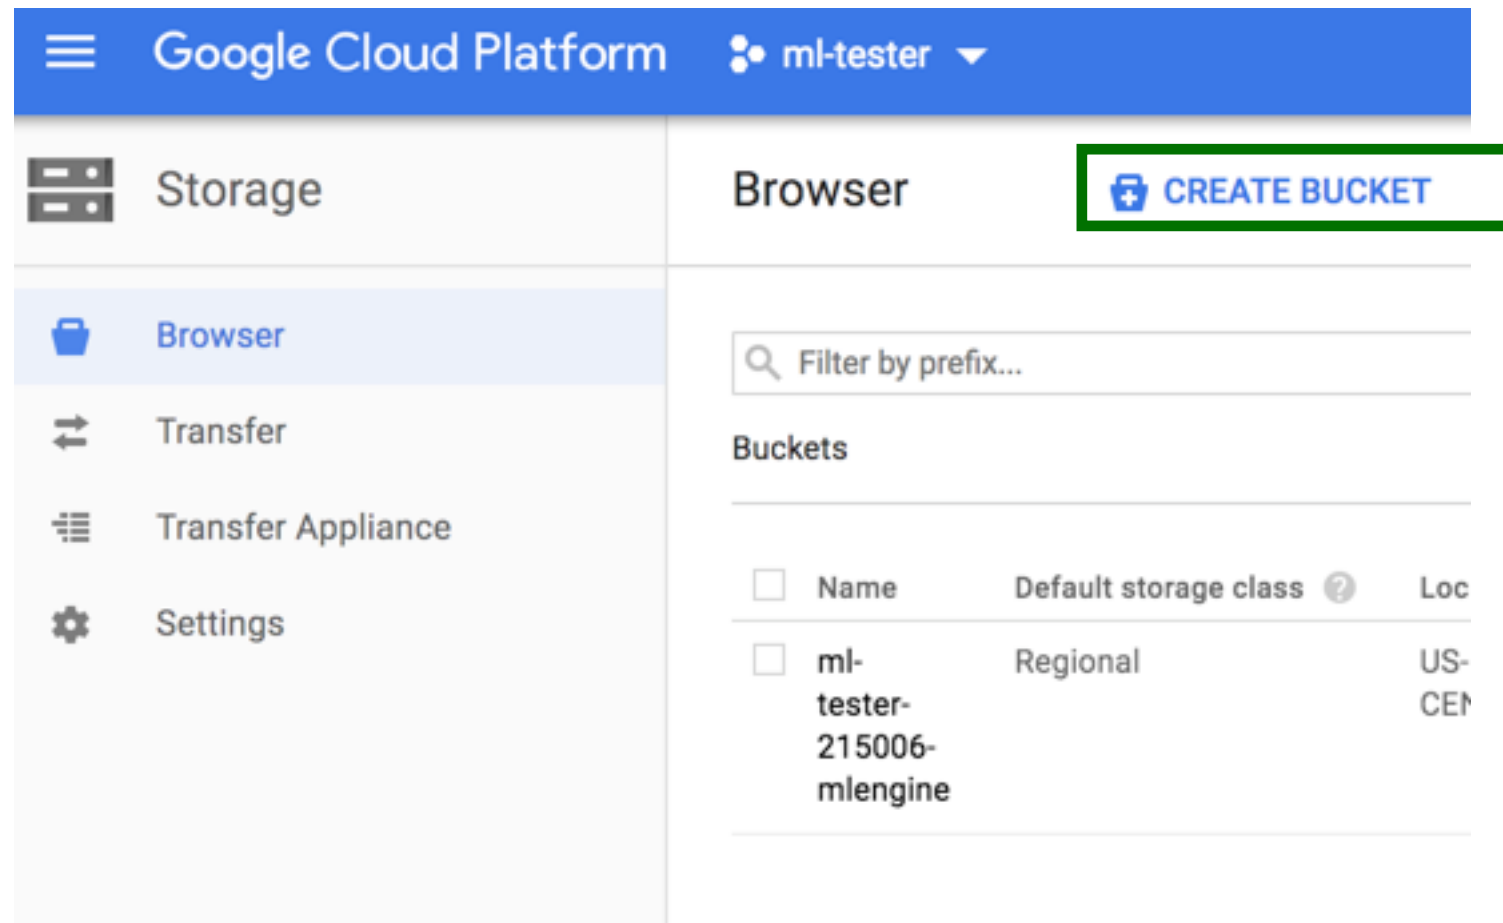
1. Select or create a GCP project.
[GO TO THE MANAGE RESOURCES PAGE](#)
2. 프로젝트에 결제가 사용 설정되어 있는지 확인하세요.
[결제 사용 설정 방법 알아보기](#)
3. Enable the Cloud Machine Learning Engine and Compute Engine APIs.
[ENABLE THE APIS](#)
4. Set up authentication:
 - a. Go to the **Create service account key** page in the GCP Console.
[GO TO THE CREATE SERVICE ACCOUNT KEY PAGE](#)
 - b. From the **Service account** drop-down list, select **New service account**.
 - c. Enter a name into the **Service account name** field.
 - d. From the **Role** drop-down list, select **Project > Owner**.
 - e. Click **Create**. A JSON file that contains your key downloads to your computer.
5. Set the environment variable `GOOGLE_APPLICATION_CREDENTIALS` to the file path of the JSON file that contains your service account key. This variable only applies to your current shell session, so if you open a new session, set the variable again.
 - Example: Linux or macOS
 - Example: Windows
6. Install and initialize the Cloud SDK.

basic flow introduction

1. create a bucket, and some folders in it
2. upload files to folders
3. create a project source tree
4. create a setup.py if you have any dependencies to install
5. create a config.yaml for your configuration of the project
6. create an empty `__init__.py` in every other folders except for the root path
7. write some codes
8. or copy some code files from local
 1. change `print()` function to logging to see a real-time logs
9. submit a job
10. watch the logs



create GCS bucket



make a bucket
@ where cloud ML engine
is available

bucket and cloud ML
trainer should be placed
@the same location

Create folder + Upload files

Upload files









Upload folder

Create folder

Delete

Filter by prefix...

[Buckets](#) / ml-tester-215006-mlengine

<input type="checkbox"/>	Name	Size	Type
<input type="checkbox"/>	 census_single_1/	—	Folder
<input type="checkbox"/>	 census_single_20/	—	Folder
<input type="checkbox"/>	 census_single_4/	—	Folder
<input type="checkbox"/>	 cifar-10-output/	—	Folder
<input type="checkbox"/>	 cifar-10/	—	Folder
<input type="checkbox"/>	 data/	—	Folder
<input type="checkbox"/>	 output/	—	Folder
<input type="checkbox"/>	 test_job6/	—	Folder

SUPER EASY

Project Structure

main application module



Create an `__init__.py` file in every subdirectory. These files are used by [Setuptools](#) to identify directories with code to package, and may be empty.

Dependencies

- runtime versions shows packages that are already installed. You need to specify which version to use.
* `--runtime_version x.x`
- standard dependencies can be defined in `setup.py` file

```
from setuptools import find_packages
from setuptools import setup

REQUIRED_PACKAGES = ['some_PyPI_package>=1.0']

setup(
    name='trainer',
    version='0.1',
    install_requires=REQUIRED_PACKAGES,
    packages=find_packages(),
    include_package_data=True,
    description='My training application package.'
)
```

find packages
of your interest
@PyPI

gcloud cli will automatically recognize the `setup.py`

configuration (1)

set option "--config" to specify the location of config.yaml

```
--config=config.yaml
```

config.yaml example

```
trainingInput:
  scaleTier: CUSTOM
  masterType: complex_model_m
  workerType: complex_model_m
  parameterServerType: large_model
  workerCount: 9
  parameterServerCount: 3
  packageUris: gs://my/trainer/path/package-0.0.0.tar.gz
  pythonModule: trainer.task
  region: us-central1
  jobDir: gs://my/training/job/directory
  runtimeVersion: 1.10
  pythonVersion: 3.5
```

see the machine types, and here

standard_gpu	A machine equivalent to standard that also includes a single NVIDIA Tesla K80 GPU. Compute Engine machine name: n1-standard-8 with one k80 GPU
--------------	---

\$1.2118 (2.4731)

standard_p100	A machine equivalent to standard that also includes a single NVIDIA Tesla P100 GPU. Compute Engine machine name: n1-standard-8 p100x1
---------------	--

\$2.6864 (5.4824)

standard_v100	A machine equivalent to a standard that also includes a single NVIDIA Tesla V100 GPU. The availability of these GPUs is in Beta launch stage. Compute Engine machine name: n1-standard-8-v100x1
---------------	--

Job directory (jobDir)

The path to a Cloud Storage location to use for job output.

configuration (2)

other common training parameters

```
gcloud ml-engine jobs submit training JOB --module-name = MODULE_NAME [--config = CONFIG]  
  [--job-dir = JOB_DIR] [--labels = [ KEY = VALUE ,...]] [--package-path = PACKAGE_PATH]  
  [--packages = [ PACKAGE ,...]] [--python-version = PYTHON_VERSION] [--region = REGION]  
  [--runtime-version = RUNTIME_VERSION] [--scale-tier = SCALE_TIER]  
  [--staging-bucket = STAGING_BUCKET] [--async | --stream-logs] [G_CLOUD_WIDE_FLAG ...]  
  [- USER_ARGS ...]
```

further description for each parameter

read/write file from GCS?

from tensorflow.python.lib.io **import** file_io

python standard library can't recognize GCS location,
but, tensorflow provides file_io packages for it

```
bucket_name = 'ml-tester-215006-mlengine'  
path = 'cifar-10'  
filename = 'cifar10_preprocess_batch_' + str(batch_id) + '.p'  
full_path = 'gs://' + bucket_name + '/' + path + '/' + filename
```

GCS location can be recognized by three parts,
“gs://” protocol, “bucket_name”, and “file_path”

```
from tensorflow.python.lib.io import file_io  
  
with file_io.FileIO('gs://.....', mode='w+') as f:  
    cPickle.dump(self.words, f)
```

Or you can read pickle file in like this:

```
file_stream = file_io.FileIO(train_file, mode='r')  
x_train, y_train, x_test, y_test = pickle.load(file_stream)
```

logging

- standard python's print function shows up "AFTER" everything is done. This is a known buffering issue.
- If you want to check a real time logs, you need to use "logging" package instead.

```
import logging

logger = logging.getLogger('simple_example')
logger.debug('debug message')
logger.info('info message')
logger.warn('warn message')
logger.error('error message')
logger.critical('critical message')
```

submitting a job (1)

with common parameters in option flags

```
gcloud ml-engine jobs submit training $JOB_NAME \
  --package-path $TRAINER_PACKAGE_PATH \
  --module-name $MAIN_TRAINER_MODULE \
  --job-dir $JOB_DIR \
  --region $REGION \
  --config config.yaml \
  -- \
  --user_first_arg=first_arg_value \
  --user_second_arg=second_arg_value
```

```
gcloud ml-engine jobs submit training $JOB_NAME \
  --scale-tier basic \
  --package-path $TRAINER_PACKAGE_PATH \
  --module-name $MAIN_TRAINER_MODULE \
  --job-dir $JOB_DIR \
  --region $REGION \
  -- \
  --user_first_arg=first_arg_value \
  --user_second_arg=second_arg_value
```

```
TRAINER_PACKAGE_PATH="/path/to/your/application/sources"
now=$(date +"%Y%m%d_%H%M%S")
JOB_NAME="your_name_$now"
MAIN_TRAINER_MODULE="trainer.task"
JOB_DIR="gs://your/chosen/job/output/path"
PACKAGE_STAGING_PATH="gs://your/chosen/staging/path"
REGION="us-east1"
RUNTIME_VERSION="1.10"
```

submitting a job (2)

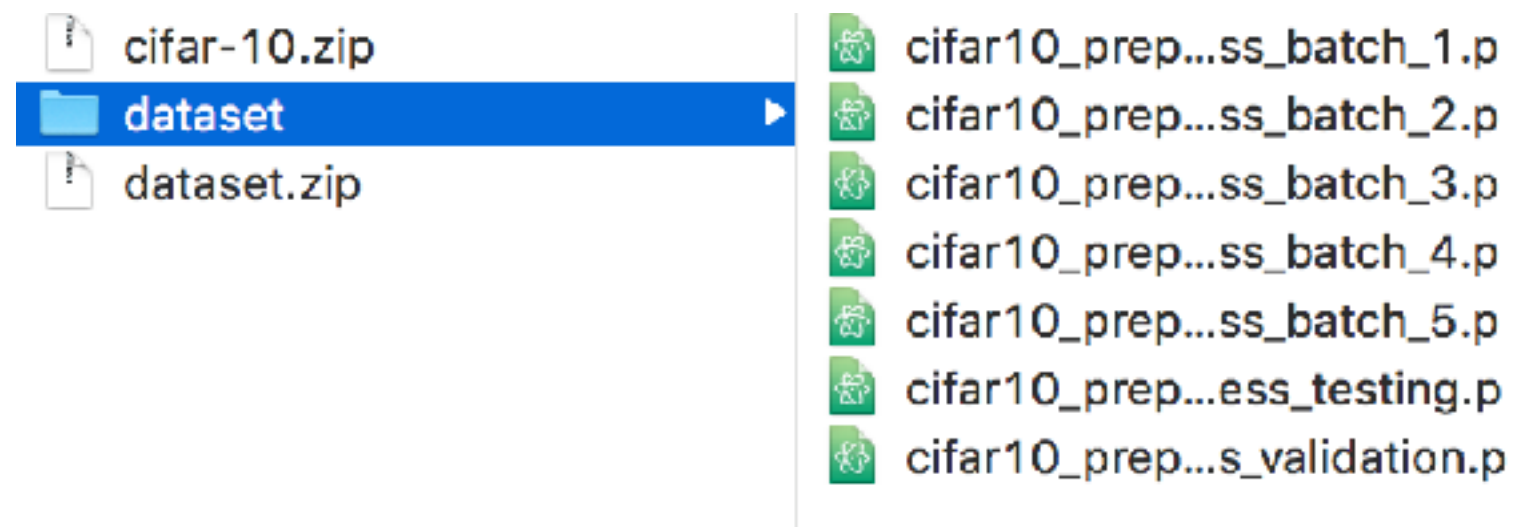
with only config.yaml configuration file,
in which every parameters are defined

```
gcloud ml-engine jobs submit training $JOB_NAME \  
  --config config.yaml \  
  -- \  
  --user_first_arg=first_arg_value \  
  --user_second_arg=second_arg_value
```

Example

- showing a simple example of CIFAR10 image classification with GoogLeNet (InceptionV1) model.
- This example will use TESLA V100/P100
- All the materials needed for this example can be downloaded from [HERE](#)

Step1 - extract materials



**cifar-10 file contains
source codes to train**

**dataset folder contains dataset to train on.
these files should be uploaded into the GCP bucket**

Step2 - upload data to GCP

Also possible to download dataset from source code and save it to the bucket

here, I will show how to upload manually though

1

2

3

4

locations supporting GPUs/TPU

	US	EUROPE	ASIA PACIFIC
Region	us-west1	us-central1	us-east1
NVIDIA TESLA K80		✓	✓
NVIDIA TESLA P100	✓	✓	✓
NVIDIA TESLA V100	✓	✓	
TPU_V2		✓	

	US	EUROPE	ASIA PACIFIC
Region	asia-south1	asia-southeast1	asia-east1
NVIDIA TESLA K80			✓
NVIDIA TESLA P100			✓
NVIDIA TESLA V100			✓
TPU_V2			

Step2 - upload data to GCP(2)

cifar10-bucket

[Objects](#) [Overview](#) [Permissions](#)

1

[Upload files](#) [Upload folder](#) [Create folder](#) [Delete](#)

[Buckets](#) / [cifar10-bucket](#)

create folder



Create folder

Name

You will create a folder named dataset in cifar10-bucket/

[CANCEL](#) [CREATE](#)

2

name folder



cifar10-bucket


[Objects](#) [Overview](#) [Permissions](#)

4

[Upload files](#) [Upload folder](#) [Create folder](#) [Delete](#)

[Buckets](#) / [cifar10-bucket](#) / [dataset](#)

upload dataset files into the folder



<input type="checkbox"/>	Name	Size	Type	Storage class	Last modified	Public access
<input type="checkbox"/>	dataset/	—	Folder	—	—	Per object

3

**check if folder has been created successfully,
and click the folder**


Step2 - upload data to GCP(3)

cifar10-bucket

[Objects](#) [Overview](#) [Permissions](#)

[Upload files](#) [Upload folder](#) [Create folder](#) [Delete](#)

[Buckets](#) / [cifar10-bucket](#) / [dataset](#)

<input type="checkbox"/>	Name	Size	Type	Storage class	Last modified	Public access ?	Encryption ?	
<input type="checkbox"/>	 cifar10_preprocess_batch_1.p	38.7 MB	application/octet-stream	Regional	9/10/18, 5:36 PM	Not public	Google-managed key	⋮
<input type="checkbox"/>	 cifar10_preprocess_batch_2.p	38.71 MB	application/octet-stream	Regional	9/10/18, 5:36 PM	Not public	Google-managed key	⋮
<input type="checkbox"/>	 cifar10_preprocess_batch_3.p	38.56 MB	application/octet-stream	Regional	9/10/18, 5:36 PM	Not public	Google-managed key	⋮
<input type="checkbox"/>	 cifar10_preprocess_batch_4.p	38.68 MB	application/octet-stream	Regional	9/10/18, 5:36 PM	Not public	Google-managed key	⋮
<input type="checkbox"/>	 cifar10_preprocess_batch_5.p	38.69 MB	application/octet-stream	Regional	9/10/18, 5:35 PM	Not public	Google-managed key	⋮
<input type="checkbox"/>	 cifar10_preprocess_testing.p	43.12 MB	application/octet-stream	Regional	9/10/18, 5:35 PM	Not public	Google-managed key	⋮
<input type="checkbox"/>	 cifar10_preprocess_validation.p	21.61 MB	application/octet-stream	Regional	9/10/18, 5:34 PM	Not public	Google-managed key	⋮

create “ckpt” folder just like “dataset” folder
- “ckpt” folder will store training checkpoint files


Step3 - upload source files

1



activate cloud shell

2



upload cifar-10.zip



check if the file exists

* uploading files will store files only under the home directory

step4-extract source files

```
~ $ mkdir ml-test  
~ $ mv cifar-10.zip ./ml-test  
~ $ cd ml-test/  
~/ml-test $ unzip cifar-10.zip
```

create a folder named “ml-test”

move cifar-10.zip file to the created folder

unzip the cifar-10.zip file after moving into the directory

```
~/ml-test $ find .  
.  
./cifar-10.zip  
./cifar-10  
./cifar-10/trainer  
./cifar-10/trainer/__init__.py  
./cifar-10/trainer/clftrainer.py  
./cifar-10/trainer/test.py  
./cifar-10/trainer/googlenet.py  
./cifar-10/trainer/cifar10_dataset.py  
./cifar-10/trainer/imgclfmodel.py  
./cifar-10/trainer/dataset.py  
./cifar-10/setup.py  
./cifar-10/config.yaml  
~/ml-test $
```

trainer class (can execute training)

where the main function is

GoogLeNet model definition

CIFAR10 dataset class (knows how to load)

base classes for model and dataset

config/setup files for cloud ml engine

step5-set some variables

```
BUCKET_NAME=cifar10-bucket  
REGION=us-central1
```

```
OUTPUT_PATH=gs://$BUCKET_NAME/output  
JOB_NAME=hello_cifar10_v100
```

```
def get_training_batches_from_preprocessed(self, batch_id, batch_size, scale_to_imagenet=False):  
    bucket_name = 'cifar10-bucket'  
    path = 'dataset'  
    filename = 'cifar10_preprocess_batch_' + str(batch_id) + '.p'  
    full_path = 'gs://' + bucket_name + '/' + path + '/' + filename  
  
    features, labels = pickle.load(file_io.FileIO(full_path, 'r'))  
  
    if scale_to_imagenet:  
        features = self.convert_to_imagenet_size(features)  
  
    return self.get_batches_from(features, labels, batch_size)  
  
def get_valid_set(self, scale_to_imagenet=False):  
    bucket_name = 'cifar10-bucket'  
    path = 'dataset'  
    filename = 'cifar10_preprocess_validation.p'  
    full_path = 'gs://' + bucket_name + '/' + path + '/' + filename  
  
    valid_features, valid_labels = pickle.load(file_io.FileIO(full_path, 'r'))  
  
    if scale_to_imagenet:  
        valid_features = self.convert_to_imagenet_size(valid_features)  
  
    return valid_features, valid_labels
```

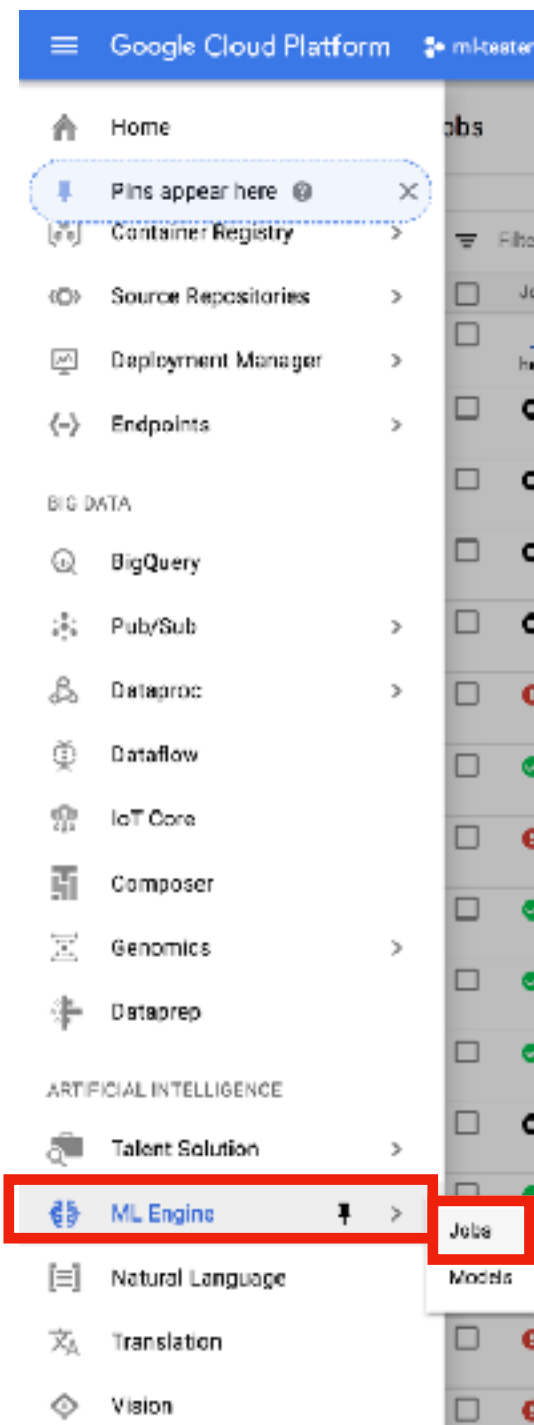
step5-set some variables(2)

```
def main():  
  
    learning_rate = 0.0001  
    epochs = 3  
    batch_size = 64  
  
    cifar10_dataset = Cifar10()  
    model = GoogLeNet()  
    trainer = ClfTrainer(model, cifar10_dataset)  
  
    trainer.run_training(epochs, batch_size, learning_rate, 'gs://cifar10-bucket/ckpt/inceptionv1.ckpt')
```

step6-run training

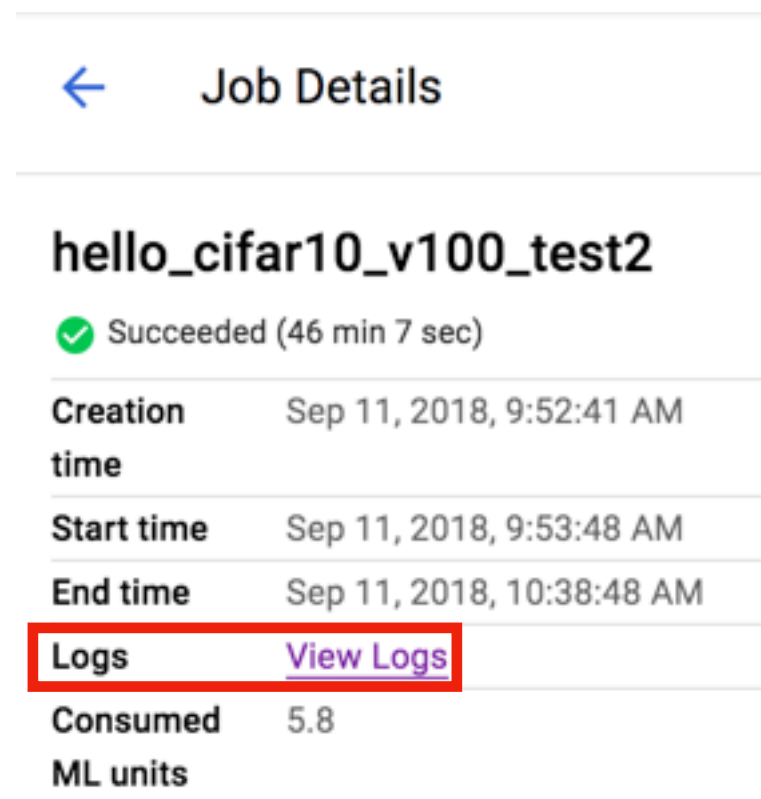
```
ml-test/cifar-10 $ gcloud ml-engine jobs submit training $JOB_NAME \  
    --package-path trainer \  
    --module-name trainer.test \  
    --job-dir $OUTPUT_PATH \  
    --region $REGION \  
    --config config.yaml \  
    --runtime-version 1.9
```


step7-watch/monitor



This screenshot shows the Google Cloud Platform ML Engine Jobs page. The 'Jobs' tab is selected. A table lists the jobs, with the job 'hello_cifar10_v100_test2' highlighted by a red box. The table has columns for Job ID, Type, Create time, Elapsed time, Logs, and Labels.

Job ID	Type	Create time	Elapsed time	Logs	Labels
hello_cifar10_v100_test2	Training	Sep 11, 2018, 9:52:41 AM	45 min 50 sec	View Logs	



step7-watch/monitor(2)

```
master-replica-0 name: Tesla V100-SXM2-16GB major: 7 minor: 9 memoryClockRate(GHz): 1.53
```

```
master-replica-0 +-----+
master-replica-0 | NVIDIA-SMI 384.111 driver version: 384.111 |
master-replica-0 |-----+-----+
master-replica-0 | GPU Name Persistence-M| Bus-Id Disp.A | Volatile Uncorr. ECC |
master-replica-0 | Fan Temp Perf Pwr:Usage/Cap| Memory-Usage | GPU-Util Compute M. |
master-replica-0 |=====+=====+=====+
master-replica-0 | 0 Tesla V100-SXM2... Off | 00000000:00:04:0 Off | 0 |
master-replica-0 | N/A 35C P0 36W / 300W | 15370MiB / 16152MiB | 0% Default |
master-replica-0 +-----+-----+-----+-----+-----+-----+
```

text epoch x text validation x

Cloud ML Job, hello_cifar10_v100_test2

All logs

Any log level

Last hour

Jump to now

Showing logs from the last hour ending at 10:41 AM (JST)

2018-09-11 10:03:03.458 JST	master-replica-0	Validation Accuracy 0.360916
2018-09-11 10:04:49.389 JST	master-replica-0	Epoch 1, Cifar-10 Batch 4: Avg. Loss: 1.69417114714
2018-09-11 10:05:38.557 JST	master-replica-0	Validation Accuracy 0.385817
2018-09-11 10:07:24.456 JST	master-replica-0	Epoch 1, Cifar-10 Batch 5: Avg. Loss: 1.63898358277
2018-09-11 10:08:14.142 JST	master-replica-0	Validation Accuracy 0.394832
2018-09-11 10:08:14.143 JST	master-replica-0	epoch: 1 is saved...
2018-09-11 10:10:03.886 JST	master-replica-0	Epoch 2, Cifar-10 Batch 1: Avg. Loss: 1.57556438192
2018-09-11 10:10:53.447 JST	master-replica-0	Validation Accuracy 0.425080
2018-09-11 10:12:38.266 JST	master-replica-0	Epoch 2, Cifar-10 Batch 2: Avg. Loss: 1.53801986343
2018-09-11 10:13:28.058 JST	master-replica-0	Validation Accuracy 0.416567
2018-09-11 10:15:12.338 JST	master-replica-0	Epoch 2, Cifar-10 Batch 3: Avg. Loss: 1.45232008704
2018-09-11 10:16:01.482 JST	master-replica-0	Validation Accuracy 0.485176
2018-09-11 10:17:46.282 JST	master-replica-0	Epoch 2, Cifar-10 Batch 4: Avg. Loss: 1.42776431175
2018-09-11 10:18:36.171 JST	master-replica-0	Validation Accuracy 0.497396
2018-09-11 10:20:20.820 JST	master-replica-0	Epoch 2, Cifar-10 Batch 5: Avg. Loss: 1.39621625079
2018-09-11 10:21:10.171 JST	master-replica-0	Validation Accuracy 0.503606
2018-09-11 10:21:10.171 JST	master-replica-0	epoch: 2 is saved...
2018-09-11 10:22:59.562 JST	master-replica-0	Epoch 3, Cifar-10 Batch 1: Avg. Loss: 1.32890029595
2018-09-11 10:23:49.459 JST	master-replica-0	Validation Accuracy 0.486979
2018-09-11 10:25:35.505 JST	master-replica-0	Epoch 3, Cifar-10 Batch 2: Avg. Loss: 1.31552324312
2018-09-11 10:26:26.220 JST	master-replica-0	Validation Accuracy 0.535056
2018-09-11 10:28:11.454 JST	master-replica-0	Epoch 3, Cifar-10 Batch 3: Avg. Loss: 1.22831901295
2018-09-11 10:29:00.640 JST	master-replica-0	Validation Accuracy 0.564704
2018-09-11 10:30:44.259 JST	master-replica-0	Epoch 3, Cifar-10 Batch 4: Avg. Loss: 1.22302149416
2018-09-11 10:31:33.785 JST	master-replica-0	Validation Accuracy 0.582131
2018-09-11 10:33:18.407 JST	master-replica-0	Epoch 3, Cifar-10 Batch 5: Avg. Loss: 1.19839818216
2018-09-11 10:34:07.951 JST	master-replica-0	Validation Accuracy 0.581731
2018-09-11 10:34:07.951 JST	master-replica-0	epoch: 3 is saved...

appendix

- how to get current project ID?
`$(gcloud config list project --format "value(core.project)")`
- with `tf.device('/device:GPU:0')`:



**Google
Cloud Platform**

Cloud Platform



**Cloud
ML Engine**

ML Engine